

Class 11: Structural Bioinformatics pt2

Mai Tamura (PID: A18594079)

Table of contents

AlphaFold DB	1
Generating Your own structure predictions	2
Custom Analysis of resulting models in R	5
Residue conservation from alignment file	8

AlphaFold DB

The EBI maintains the largest database of AlphaFold structure prediction models at:
<https://alphafold.ebi.ac.uk>

From last class (before Halloween), we saw that the PDB had 244,290 (Oct 2025)

The total number of protein sequences in UniProtKB is 199,579,901

Key point: This is a fraction of sequence space that has structural coverage (0.12%)

```
244290/199579901 * 100
```


```
[1] 0.1224021
```

AFDB is attempting to address this gap...


There are Two “Quality Scores” from AlphaFold one for residues (i.e. each amino acid) called **pLDDT score**. The other **PAE** score that measures the confidence in the relative position of two residues (i.e. a score for every pair of residues).

Generating Your own structure predictions

Figures of 5 generated HIV-PR models




At the top model



model 1



and model 5



```
unzip("HIVPR_dimer_23119.result.zip")
```

Custom Analysis of resulting models in R

Read key result files into R. The first thing I need to know is what my results directory/folder is called (i.e. its name is different from)

```
results_dir <- "HIVPR_dimer_23119"  
pdb_files <- list.files(path=results_dir,  
                         pattern="*.pdb",  
                         full.names = TRUE)  
basename(pdb_files)
```

```
[1] "HIVPR_dimer_23119_unrelaxed_rank_001_alphaFold2_multimer_v3_model_4_seed_000.pdb"  
[2] "HIVPR_dimer_23119_unrelaxed_rank_002_alphaFold2_multimer_v3_model_1_seed_000.pdb"  
[3] "HIVPR_dimer_23119_unrelaxed_rank_003_alphaFold2_multimer_v3_model_5_seed_000.pdb"
```

```
[4] "HIVPR_dimer_23119_unrelaxed_rank_004_alphaFold2_multimer_v3_model_2_seed_000.pdb"
[5] "HIVPR_dimer_23119_unrelaxed_rank_005_alphaFold2_multimer_v3_model_3_seed_000.pdb"
```

```
library(bio3d)

pdbs <- pdbaln(pdbs_files, fit=TRUE, exefile="msa")
```

Reading PDB files:

```
HIVPR_dimer_23119/HIVPR_dimer_23119_unrelaxed_rank_001_alphaFold2_multimer_v3_model_4_seed_000.pdb
HIVPR_dimer_23119/HIVPR_dimer_23119_unrelaxed_rank_002_alphaFold2_multimer_v3_model_1_seed_000.pdb
HIVPR_dimer_23119/HIVPR_dimer_23119_unrelaxed_rank_003_alphaFold2_multimer_v3_model_5_seed_000.pdb
HIVPR_dimer_23119/HIVPR_dimer_23119_unrelaxed_rank_004_alphaFold2_multimer_v3_model_2_seed_000.pdb
HIVPR_dimer_23119/HIVPR_dimer_23119_unrelaxed_rank_005_alphaFold2_multimer_v3_model_3_seed_000.pdb
....
```

Extracting sequences

```
pdbs/seq: 1 name: HIVPR_dimer_23119/HIVPR_dimer_23119_unrelaxed_rank_001_alphaFold2_multimer_v3_model_4_seed_000.pdb
pdbs/seq: 2 name: HIVPR_dimer_23119/HIVPR_dimer_23119_unrelaxed_rank_002_alphaFold2_multimer_v3_model_1_seed_000.pdb
pdbs/seq: 3 name: HIVPR_dimer_23119/HIVPR_dimer_23119_unrelaxed_rank_003_alphaFold2_multimer_v3_model_5_seed_000.pdb
pdbs/seq: 4 name: HIVPR_dimer_23119/HIVPR_dimer_23119_unrelaxed_rank_004_alphaFold2_multimer_v3_model_2_seed_000.pdb
pdbs/seq: 5 name: HIVPR_dimer_23119/HIVPR_dimer_23119_unrelaxed_rank_005_alphaFold2_multimer_v3_model_3_seed_000.pdb
```

```
pdbs
```

[Truncated_Name:1]	HIVPR_dime	1	.	.	.	50
		PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWPKPMIGGI				
[Truncated_Name:2]	HIVPR_dime	PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWPKPMIGGI				
[Truncated_Name:3]	HIVPR_dime	PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWPKPMIGGI				
[Truncated_Name:4]	HIVPR_dime	PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWPKPMIGGI				
[Truncated_Name:5]	HIVPR_dime	PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWPKPMIGGI				

		1	.	.	.	50
		51	.	.	.	100
[Truncated_Name:1]	HIVPR_dime	GGFIKVRQYDQILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNFP				
[Truncated_Name:2]	HIVPR_dime	GGFIKVRQYDQILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNFP				
[Truncated_Name:3]	HIVPR_dime	GGFIKVRQYDQILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNFP				
[Truncated_Name:4]	HIVPR_dime	GGFIKVRQYDQILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNFP				
[Truncated_Name:5]	HIVPR_dime	GGFIKVRQYDQILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNFP				

```

51 . . . .
101 . . . .
150

[Truncated_Name:1] HIVPR_dime QITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWPCKMIGGIG
[Truncated_Name:2] HIVPR_dime QITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWPCKMIGGIG
[Truncated_Name:3] HIVPR_dime QITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWPCKMIGGIG
[Truncated_Name:4] HIVPR_dime QITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWPCKMIGGIG
[Truncated_Name:5] HIVPR_dime QITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWPCKMIGGIG
*****  

101 . . . .
150

151 . . .
198

[Truncated_Name:1] HIVPR_dime GFIKVRQYDQILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNF
[Truncated_Name:2] HIVPR_dime GFIKVRQYDQILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNF
[Truncated_Name:3] HIVPR_dime GFIKVRQYDQILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNF
[Truncated_Name:4] HIVPR_dime GFIKVRQYDQILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNF
[Truncated_Name:5] HIVPR_dime GFIKVRQYDQILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNF
*****  

151 . . .
198

Call:  

  pdbaln(files = pdb_files, fit = TRUE, exefile = "msa")  

Class:  

  pdbs, fasta  

Alignment dimensions:  

  5 sequence rows; 198 position columns (198 non-gap, 0 gap)  

+ attr: xyz, resno, b, chain, id, ali, resid, sse, call  

  pdb <- read.pdb("1hsg")

```

Note: Accessing on-line PDB file

```
m1 <- read.pdb(pdb_files[1])  
m1
```

```
Call: read.pdb(file = pdb_files[1])
```

```
Total Models#: 1
Total Atoms#: 1514, XYZs#: 4542 Chains#: 2 (values: A B)
```

```
Protein Atoms#: 1514 (residues/Calpha atoms#: 198)
Nucleic acid Atoms#: 0 (residues/phosphate atoms#: 0)
```


```
Non-protein/nucleic Atoms#: 0 (residues: 0)
Non-protein/nucleic resid values: [ none ]
```

Protein sequence:

```
PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWPKMIGGIGGF IKVRQYD
QILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNFPQITLWQRPLVTIKIGGQLKE
ALLDTGADDTVLEEMSLPGRWPKMIGGIGGF IKVRQYDQILIEICGHKAIGTVLVGPTP
VNIIGRNLLTQIGCTLNF
```

+ attr: atom, xyz, calpha, call

```
plot.bio3d(m1$atom$b[m1$calpha], typ="l", ylim=c(0,100))
```



Residue conservation from alignment file

Find the large AlphaFold alignment file

```
aln_file <- list.files(path=results_dir,
                        pattern=".a3m$",
                        full.names = TRUE)
aln_file
```

```
[1] "HIVPR_dimer_23119/HIVPR_dimer_23119.a3m"
```

Read this into R

```
aln <- read.fasta(aln_file[1], to.upper = TRUE)
```

```
[1] " ** Duplicated sequence id's: 101 **"
[2] " ** Duplicated sequence id's: 101 **"
```

How many sequences are in this alignment

```
dim(aln$ali)
```

```
[1] 5397 132
```

We can score residue conservation in the alignment with the `conserv()` function.

```
sim <- conserv(aln)

plotb3(sim[1:99], sse=trim.pdb(pdb, chain="A"),
       ylab="Conservation Score")
```

